



Toy Model — Ambiguity: interpolation between features 1 and 2
($|\cosine| = 0.810$; $magnitude=0.5$)


SAE: per-feature activations



SAE: credit allocation
(entropy at midpoint: -0.000 / 1.609)



VAE: $|\mu|$ per feature ($\pm\sigma$ shaded for fi, fj)



VAE: posterior sigma
(entropy at midpoint: 0.376 / 1.609)

